Continuous Top-k Spatial-Keyword Search on Dynamic Objects

Yuyang Dong[†]   Hanxiong Chen[‡]   Chuan Xiao[§]

Kunihiro Takeoka[†]   Masafumi Oyamada [†]   Jeffrey Xu Yu[♮]   Hiroyuki Kitagawa[‡]

[†]*NEC Corporation, Japan*   [‡]*University of Tsukuba, Japan*
[§]*Osaka University & Nagoya University, Japan*   [♮]*The Chinese University of Hong Kong, China*
{y-dong@aj, k-takeoka@az, m-oyamada@cq}.jp.nec.com
{chx, kitagawa}@cs.tsukuba.ac.jp   chuanx@ist.osaka-u.ac.jp   yu@se.cuhk.edu.hk

ABSTRACT

As the popularity of SNS and GPS-equipped mobile devices rapidly grows, numerous location-based applications have emerged. A common scenario is that a large number of users change location and interests from time to time; e.g., a user watches news, videos, and blogs while moving outside. Many online services have been developed based on continuously querying spatial-keyword objects. For instance, a real-time coupon delivery system searches for potential customers using their locations and interested keywords, and sends coupons to attract potential customers.

In this paper, we investigate the case of *dynamic* spatial-keyword objects whose locations and keywords change over time. We study the problem of continuously tracking top-k dynamic spatial-keyword objects for a given set of queries. Answering this type of queries benefit many location-aware services such as E-coupon recommender systems, drone delivery, and self-driving stores. We develop a solution based on a grid index. To deal with the changing locations and keywords of objects, our solution first finds the set of queries whose results are affected by the change and then updates the results of these queries. We propose a series of indexing and query processing techniques to accelerate the two procedures. We also discuss batch processing to cope with the case when multiple objects change locations and keywords in a time interval. Experimental results on real and synthetic datasets demonstrate the efficiency of our method and its superiority over alternative solutions.
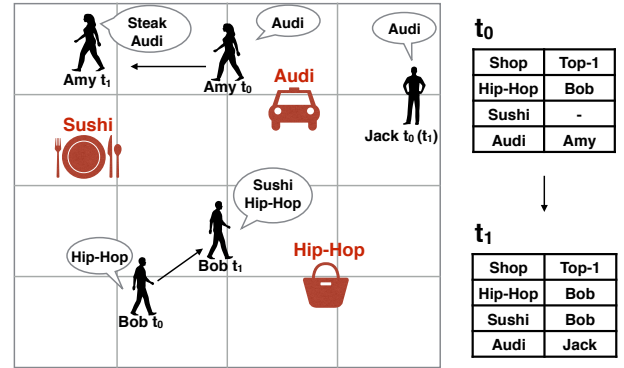
1. INTRODUCTION

Figure 1: E-coupon recommender system.

The prevalence of GPS-equipped mobile devices generates a massive volume of spatial-keyword data. Processing spatial-keyword data is an essential procedure in location-aware applications such as traffic monitoring, unsafe moving object detection, and recommender systems for mobile users. Moreover, many innovative applications in the upcoming 5G network era also rely on continuously querying spatial-keyword data. E.g., for drone delivery [2], we can monitor drones' positions and the information of packages to confirm they are working properly; for self-driving grocery stores [20], users can view the locations of stores and the goods on sale before they order.

Despite various types of queries on spatial-keyword data being studied in the last decade, existing studies focus on dealing with static objects or moving objects that only change locations [13, 14, 17, 18, 26, 27, 29, 33, 34]. Yet, real-world objects are often *dynamic* and change both locations and keywords over time. In this paper, we explore the scenario of processing top-k dynamic spatial-keyword objects: given a set of dynamic spatial-keyword objects whose locations and keywords may vary from time to time, and a set of (static) queries represented by locations and keywords, our task is to monitor the top-k results for each query, ranked by a scoring function of a weighted average of spatial and keyword similarities. Solving this problem benefits many underlying applications, e.g., location-aware recommender systems and drone/self-driving store monitoring. We show a motivating example as follows.

Consider an E-coupon recommender system in Figure 1, where a hip-hop cloth store, a sushi restaurant and an Audi

dealer are registered as queries. Three users (objects) are using mobile phones. Suppose we search for top-1 results; i.e, the system monitors the top-1 users w.r.t. the three stores and sends out coupons. At time t_0, Bob is watching a hip-hop music video, Amy is searching for Audi, and Jack is reading the news on Audi. Bob becomes the top-1 result of the hip-hop store, as he is the only user that matches keyword hip-hop. Both Amy and Jack are associated with keyword Audi, but Amy is closer to the Audi dealer and becomes the top-1 result. Nobody is associated with keyword sushi. So the top-1 result for the sushi restaurant is empty. At time t_1, Bob is moving northeast and starts to watch an eating show about sushi. So his keywords become hip-hop and sushi [1]. Amy moves away from the Audi dealer and searches "steak near me". So her keywords become Audi and steak. Jack is staying and still reading the news on Audi. The system keeps Bob at the top-1 result of the hip-hop store and recognizes him as the new top-1 of the sushi restaurant. As Jack becomes closer than Amy to the Audi dealer, he replaces Amy as the top-1 of the Audi dealer.

To the best of our knowledge, this is the first work targeting the case when objects change both spatial and textual attributes [2]. For real-time applications, the key issue is to efficiently update top-k results for the queries whenever an object updates its status. Although one may regard the update of an object as a deletion followed by an insertion, thus to convert dynamic objects to static ones, simply applying existing methods for related problems – CIQ [5] and SKYPE [24] for location-aware publish/subscribe systems – does not deliver sufficient efficient query processing. Our experiments show that adaptation of these methods spend 50ms to 150ms to process a single object update on datasets of around 200K queries, hence unlikely to keep up with the very frequent update in real-time applications. Next we summarize the challenges and our method.

1.1 Challenges and Contributions

The first challenge originates from the large number of queries. It is prohibitive to check every query and recompute top-k results. On the other hand, the number of **affected queries** (i.e., the queries that change their top-k results) is usually very small once an object update occurs. This challenge is akin to the case in location-aware publish/subscribe systems. Both CIQ [5] and SKYPE [24] utilize an inverted index for keywords and a quadtree for spatial information to quickly identify the queries (referred to as subscriptions in [5, 24]) affected by an object update (referred to as messages in [5,24]). SKYPE further adopts the prefix filtering technique used in set similarity search [3, 28] to limit the inverted index access to only keywords with high weights. However, whenever an object update occurs, both CIQ and SKYPE have to traverse the quadtree and compute score upperbounds to check whether they need to process the queries in a cell, hence resulting in considerable cell access. In addition, both methods index single keywords in the inverted index, which turns out to be inefficient for our problem. This is because it is a common scenario that an object has only frequent keywords after an update, thereby resulting in access to long postings lists in the inverted index and considerable overhead.

To address the first challenge, we index objects and queries in a grid index and propose the notions of cell-cell links (*CC links*) and *l-signatures*. CC links model the cells in a grid as a directed graph. Whenever an object update occurs, through the outgoing CC links of the cell of the new location, we can quickly identify the cells that need look-up and avoid accessing the cells that can be pruned. *l-signatures* are combinations of l keywords. Despite the existence of frequent keywords, the frequencies of their combinations are often significantly smaller. By indexing *l-signatures* of queries and selecting a good set of *l-signatures* from the updated object, the cost of inverted index access can be remarkably reduced.

The second challenge is to compute the top-k results of the affected queries. Since the update of an object may cause the object to move out of the top-k results of some queries (e.g., when an object moves away from a query), it is time-consuming to **refill the top-k** of these queries from scratch. Most related studies adopt a buffering strategy to store a list of non-top-k objects for each query. The buffered objects are used when a top-k refilling is needed. E.g., the kmax buffering [31] maintains top-k' results where k' is a value between k and a maximum buffer size k_{max}; SKYPE uses a cost-based k-skyband buffer by setting a threshold of the score. Such object-based buffering is inefficient for our problem because a reordering of the objects in the buffer is required whenever an update occurs.

To address the second challenge, we first propose a *bounding* method to leverage the upper and lower bounds of scores of the objects in a cell. The key observation is that the score bounds of the objects in a cell do not change frequently. To improve the performance, we also observe that the new top-k result of a query is either the updated object or the $(k + 1)$-th result prior to the update. Thus, on top of the bounding method, we propose an enhanced method that maintains a *list of cells* for each query such that $(k+1)$-th object of the query is guaranteed to reside in one of these cells. We derive the condition of these cells and address the issue of the list update.

We extend our method to batch processing, i.e., multiple objects change status in an time interval. By taking into account of sharing computation for multiple objects, the technique can be used to handle the case when top-k results are supposed to be reported periodically. In addition, the initialization of top-k results and the update of queries are discussed. We conduct experiments on real and synthetic datasets. The results demonstrate that the proposed techniques are effective in reducing query processing time and contribute to a substantial overall speed-up of 5 to 20 times over alternative solutions, hence reducing the average query processing time to milliseconds.

Our contributions are summarized as follows.

- We study a new type of queries to continuously search top-k results for dynamic spatial-keyword objects that may change locations and keywords over time.
- We propose a query processing method that comprises an affected query finder and a top-k refiller to address the technical challenges of the studied problem. We devise a series of data structures (CC links, etc.) for the two components.
- We extend our method to handle the case of batch processing.
- We conduct extensive experiments. The results demonstrate the effectiveness of the components of our method and the superiority of our method over alternative solutions in speed.

The rest of the paper is organized as follows. Section 2 reviews related work. Section 3 introduces preliminaries and defines the problem. Section 4 provides the overview of our solution. Sections 5 and 6 present the algorithms of finding affect queries finder and refilling top-k. The extension to batch processing is covered by Section 7. The initialization and the update of

---

[1] We assume that previous keywords do not immediately disappear in this example. The setting depends on the application.
[2] Compared to the preliminary version of this paper [10], we made substantial improvement in our methods of affected query finder and top-k refiller, along with more experiments conducted.

Table 1: Comparison to related studies.

Research	Objects	Queries	Attributes
This work	dynamic	static	spatial-keyword
[5, 24]	streaming	streaming	spatial-keyword
[13, 14, 26, 27, 34]	static	moving	spatial-keyword
[6, 8, 16]	static	static	spatial-keyword
[17, 18, 29, 33]	moving	static	spatial

Table 2: Frequently used notations.

Symbol	Description
o, O	an object, a set of objects
q, Q	a query, a set of queries
$SimST(o, q)$	the score of o w.r.t. q
t	a timestamp
$o^t, o^{t'}$	the status of an object o at time t/t'
$q.obj(k, t)$	the k-th object of q (ranked by score) at t
$q.obj(1..k, t)$	the set of top-k objects of q at time t
$q.score(k, t)$	the score of $q.obj(k, t)$
c, C	a cell and a grid index
$c^t, c^{t'}$	the status of a cell c at time t/t'
$o.c, q.c$	the cell in which o/q is located
Q_{prev}	the queries of which o is a top-k result at t
Q_{next}	the queries s.t. $SimST(o^{t'}, q) > q.score(k, t)$
τ	a keyword similarity threshold
$c^t.\psi_{\max}, c^t.\psi_{\min}$	the keywords in c^t, weighted by max/min in c^t
$maxscore(c^t, q)$	max. score of the objects in c to q at t
$minscore(c^t, q)$	min. score of the objects in c to q at t
$maxminscore_{<k}$	max$\{ minscore(c^t, q) \}$, $c \in C$ and $maxscore(c^t, q) < q.score(k, t)$
$q.CL$	a cell list of q at t

queries are discussed in Section 8. Experimental results are reported in Section 9. Section 10 concludes the paper.

2. RELATED WORK

Location-aware publish/subscribe systems. The studies most related to our work focus on dealing with spatial-keyword objects in a publish/subscribe system. Users register their interests as continuous queries into the system, and then new streaming objects are delivered to relevant users. Keyword boolean matching was studied in [4, 15, 25], spatial range query was studied in [7], and the case of scoring function combining spatial and keyword similarities were considered in [5, 24]. The major difference in this line of work from ours is that they do not consider dynamic objects. For our problem, although the techniques in [5, 24] can be adapted to find affected queries, they cannot be used to compute top-k results for these queries, because [5] only supports insertion of objects and [24] is based on a sliding window, whose size has to be infinitely small to capture the semantics of object update in our problem.

Moving queries on static objects. While we focus on dynamic objects and static queries, a body of work studied the case of moving queries on static objects. Wu *et al.* [26] proposed to answer continuously moving top-k spatial-keyword (MkSK) queries using safe-regions on multiplicatively weighted Voronoi cells. Huang *et al.* [14] studied MkSK queries with a general weighted sum ranking function and proposed to use hyperbola-based safe-regions to filter objects. Zheng *et al.* [34] studied continuous boolean top-k spatial-keyword queries in a road network. Guo *et al.* [13] studied continuous top-k spatial-keyword queries with a combined ranking function.

Snapshot spatial-keyword search. Searching static geo-textual objects for spatial-keyword queries have been extensively studied, e.g., for boolean matching [9] or using a scoring function [21]. We refer readers to [6, 8, 16] for various problem settings and methods. The studies in this category focus on a snapshot query on static datasets, whereas our problem is based on a continuous query with dynamic spatial-keyword objects.

Monitoring moving objects. Another line of work [17, 18, 29, 33] aims at keeping the kNN moving objects w.r.t. a fixed query point. These solutions only consider the spatial similarity and thus cannot be directly used for our problem. E.g., the objects outside the influence region [17] may be high in keyword similarity and ranked higher than those inside the influence region, rendering the filtering ineffective.

Table 1 compares our work to existing studies. Streaming objects/queries appear in a sliding window but do not change attributes. Moving objects/queries only change locations.

3. PRELIMINARIES

Definition 1 (Dynamic Spatial-Keyword Object) *A dynamic spatial-keyword object o is a pair $(o.\rho, o.\psi)$. $o.\rho$ is the*

up-to-date location of o, represented by spatial coordinates. $o.\psi$ *keeps track of the up-to-date keywords of o, represented by a set. Both $o.\rho$ and $o.\psi$ are dynamic and change over time.*

Definition 2 (Spatial-Keyword Query) *A spatial-keyword query q is a triplet $(q.\rho, q.\psi, q.\alpha)$, where $q.\rho$ is a location, $q.\psi$ is a set of keywords, and $q.\alpha$ is a parameter to balance spatial and keyword similarities. The three parameters are all static.*

Example 1 *Consider the example In Figure 1. We regard users as objects and shops as queries. At time t_0, we have three objects: $(o_1.\rho, o_1.\psi)$, $(o_2.\rho, o_2.\psi)$, and $(o_3.\rho, o_3.\psi)$, for Amy, Bob, and Jack, respectively. $o_1.\psi = \{ \texttt{Audi} \}$. $o_2.\psi = \{ \texttt{hip-hop} \}$. $o_3.\psi = \{ \texttt{Audi} \}$. At time t_1, o_1 and o_2 are updated to $(o_1.\rho', o_1.\psi')$ and $(o_1.\rho', o_2.\psi')$, respectively, where $o_1.\psi' = \{ \texttt{Audi}, \texttt{steak} \}$ and $o_2.\psi' = \{ \texttt{hip-hop}, \texttt{sushi} \}$. The three shops are represented by three queries q_1, q_2, and q_3. $q_1.\psi = \{ \texttt{hip-hop} \}$. $q_2.\psi = \{ \texttt{sushi} \}$. $q_3.\psi = \{ \texttt{Audi} \}$.*

Although we assume static queries, our method can be extended to support dynamic queries, because updating a query is equivalent to deleting a previous query and inserting a new one. We leave this case to Section 8. For brevity, we call a dynamic spatial-keyword object an *object*, and a spatial-keyword query a *query*. Our scoring function [3] is defined as follows.

Definition 3 (Spatial-Keyword Similarity) *Given an object o and a query q, their spatial-keyword similarity is*

$$SimST(o, q) = q.\alpha \cdot SimS(o.\rho, q.\rho) + (1 - q.\alpha) \cdot SimT(o.\psi, q.\psi). \quad (1)$$

For simplicity, we call $SimST(o, q)$ the *score* of o w.r.t. q, and the score of o when context is clear. It is a weighted average of spatial similarity $SimS$ and keyword (textual) similarity $SimT$. $SimS$ is calculated by the normalized Euclidean similarity:

$$SimS(o.\rho, q.\rho) = 1 - \frac{Dist(o.\rho, q.\rho)}{maxDist}, \quad (2)$$

---

[3] This scoring function is also used in [24].

Figure 2: Solution overview.


Figure 3: Index structure.

where $Dist(\cdot, \cdot)$ measures the Euclidean distance and $maxDist$ is the maximum Euclidean distance in the space. $SimT$ is calculated using the keywords in $o.\psi$ and $q.\psi$:

$$SimT(o.\psi, q.\psi) = \sum_{w \in o.\psi \cap q.\psi} wt(o.\psi, w) \cdot wt(q.\psi, w), \quad (3)$$

where $wt(o.\psi, w)$ (or $wt(q.\psi, w)$) denotes the weight of keyword w in $o.\psi$ (or $q.\psi$). As such, $SimT$ exactly captures the cosine similarity between $o.\psi$ and $q.\psi$. We consider the following weighting scheme: the weights of the keywords in an object or a query are normalized to unit length, and the weight of each keyword is proportional to the keyword's inverse document frequency (idf) in a corpus. Static idf is assumed.

By the above definitions, our problem is defined as follows.

Problem 1 (Continuous Top-k Spatial-Keyword Search on Dynamic Objects) *Given a set of dynamic objects O and a set of queries Q, our goal is to monitor the top-k objects $o \in O$ ranked by descending order of $SimST(o, q)$ for each query $q \in Q$ at each timestamp.*

To ensure the results are relevant, we demand that every object in the top-k results of a query contain at least one common keyword with the query. In addition, our method can be easily extended to support the case when each query has a query-specified k, though a global k is assumed here. Figure 1 shows an example of the top-1 results for three queries at two timestamps.

Table 2 lists the notations frequently used in this paper. We use superscript t/t' to denote the status of an object, a cell in a grid, or a grid at time t/t'.

4. SOLUTION OVERVIEW

We focus on in-memory solutions to the problem. Figure 2 shows an overview of our solution. We employ a grid index. An object or a query is indexed in the cell in which it is located. We assume that in the initial state, there have already been a set of objects and a set of queries, and the top-k results of the queries have been computed. Since our focus is to solve the dynamic case of the problem, we leave the details of the initialization to Section 8. When the status of an object changes, the grid index is updated. Then we update the top-k results for the set of queries. Two modules are designed for efficient query processing: (1) an **affected query finder** that finds the set of affected queries (by "affected", we mean at least one top-k object of this query is replaced or changes its similarity w.r.t. the query), and (2) a **top-k refiller** that updates top-k results of these affected queries.

5. AFFECTED QUERY FINDER

When a new status of an object is received, the affected query finder finds the affected queries whose top-k results or their scores need to be updated. Given o^t and $o^{t'}$, the stati of an object o at two contiguous timestamps t and t', the change from o^t to $o^{t'}$ affects two sets of queries: (1) Q_{prev}: the set of queries such that o is a top-k result of q at t, and (2) Q_{next}: the set of queries such that $SimST(o^{t'}, q) > q.score(k, t)$. Q_{prev} refers to the queries to which o^t potentially moves out of top-k, while Q_{next} refers to the queries to which $o^{t'}$ is guaranteed to be a top-k result. They may overlap because o may stay as a top-k result despite the change of status. It is easy to see that the other queries in Q can be safely excluded.

To compute Q_{prev}, we assume that the top-k results of all the queries at t have been correctly computed. Hence, an *object-query map* (implemented as an inverted index) can be employed to map o to a list of queries of which o is a top-k result. The map is updated whenever a top-k result of a query changes. Computing Q_{next} is much more challenging. A sequential scan of Q is too time-consuming for real applications. We propose a method of computing Q_{next} composed of the following two procedures: (1) exploiting the grid index to find out the cells that may contain a query in Q_{next}, and (2) for each cell in (1), identifying the queries that share a necessary number of keywords with o using an efficient search algorithm. Two data structures, cell-cell link and l-signatures, are developed to handle the two procedures. Next we introduce them respectively.

5.1 Cell-Cell Link

Since the keyword similarity $SimT$ is no greater than 1 and the minimum distance between two cells is static, it is easy to derive a score upperbound for a query located in a different cell from $o^{t'}$'s. CIQ and SKYPE index queries in a quadtree and exploit this bound. By comparing the bound with the score of the k-th object at t, unpromising cells (i.e., those too far from $o^{t'}$ to have a bound over the k-th object's score) can be pruned by traversing the quadtree. When adapting CIQ or SKYPE to our problem, a caveat is that the index access could be expensive, because the pruning at the first few depths of the quadtree is less effective, meaning that we have to go deeper in the quadtree and access quite a number of cells. Moreover, even if a cell can be pruned, we only get aware of this after comparing the score bound with the smallest k-th object's score in this cell. Such comparison is invoked many times and significantly affects the query processing performance. Seeing this inefficiency, we seek a solution to finding Q_{next} by answering the following questions: can we avoid the comparison so that the unpromising cells are *not even accessed*?

Given $o^{t'}.c$, the cell in which $o^{t'}$ is located, our basic idea is

---
Algorithm 1: AQF$(o^t, o^{t'})$

Input : the stati of o at t and t'
Output : affected queries Q_{prev} and Q_{next}

1 $Q_{prev} \leftarrow$ the queries in the object-query map of o;
2 **foreach** c connected by an outgoing CC link of $o^{t'}.c$ **do**
3 $\quad \lfloor \; Q_{next} \leftarrow Q_{next} \cup$ GetQueryIn$(c, o^{t'})$;
4 **return** Q_{prev}, Q_{next}

---

to store at $o^{t'}.c$ a set of links that directly point to the cells having at least one possible query in Q_{next}. In doing so, we can quickly identify these cells and circumvent the access to the unpromising ones. We call these links *cell-cell links* and *CC links* for short. In this sense, the cells in the grid are the vertices of a directed graph, and there is an edge from $o^{t'}.c$ to a cell c if c potentially has a query in Q_{next}. For the sake of efficiency, the update to the CC links needs to be infrequent.

Based on this idea, we first derive the upperbound of score across cells. Given $o^{t'}.c$ and $q.c$ (the cell in which q is located), by Equation 1, we have an upperbound of score by considering their cells, assuming the keyword similarity is 1:

$$UB_{SimST(o^{t'}.c,q.c)} = q.\alpha \cdot UB_{SimS(o^{t'}.c,q.c)} + 1 - q.\alpha. \quad (4)$$

$UB_{SimS(o^{t'}.c,q.c)}$ is the upperbound of spatial similarity between $o^{t'}.c$ and $q.c$: $UB_{SimS(o^{t'}.c,q.c)} = 1 - \frac{minD(o^{t'}.c,q.c)}{maxDist}$, where $minD(o^{t'}.c, q.c)$ is the minimum distance from $o^{t'}.c$ to $q.c$. By comparing $UB_{SimST(o^{t'}.c,q.c)}$ with $q.score(k,t)$, we have the following lemma.

Lemma 1 *Consider an object $o^{t'}$ and a query $q \in Q_{next}$.* $UB_{SimS(o^{t'}.c,q.c)} > \frac{q.score(k,t)+q.\alpha-1}{q.\alpha}$.

PROOF. By the definition of Q_{next}, $SimST(o^{t'}, q) > q.score(k,t)$. Because $UB_{SimST(o^{t'}.c,q.c)}$ is the score upperbound of $o^{t'}$ and q, $UB_{SimST(o^{t'}.c,q.c)} > q.score(k,t)$. By Equation 4, $q.\alpha \cdot UB_{SimS(o^{t'}.c,q.c)} + 1 - q.\alpha > q.score(k,t)$. Therefore, $UB_{SimS(o^{t'}.c,q.c)} > \frac{q.score(k,t)+q.\alpha-1}{q.\alpha}$. $\square$

Intuitively, the lemma means that a cell c has a query in Q_{next}, only if it is close enough to $o^{t'}.c$. For each cell c, it is obvious that only the query yielding the minimum $\frac{q.score(k,t)+q.\alpha-1}{q.\alpha}$ has to be considered. We connected the cells that satisfy the condition in Lemma 1 by CC links, each of which is an edge from $o^{t'}.c$ to c. It can be seen that only when the minimum $\frac{q.score(k,t)+q.\alpha-1}{q.\alpha}$ of a cell changes, the incoming CC links of this cell are updated. Hence we can achieve the goal of infrequent update.

Example 2 *In Figure 3, the CC links of c_1 are displayed by solid lines. Suppose o_2 is the object that changes status. Since $o_2^{t'}$ is located in c_1, we use c_1's outgoing CC links to retrieve Q_{next}. So only c_4, c_5, and c_8 (just for illustration) are accessed.*

The pseudo-code of our affected query finder is given in Algorithm 1. Q_{prev} is obtained by the aforementioned object-query map (Line 1). For Q_{next}, we utilize the CC links from $o^{t'}.c$ to find the cells that may have a query in Q_{next} (Line 2), each of which processed by a function GetQueryIn (Line 3), which will be introduced next.

5.2 l-Signatures

Recall that in Equation 4, the keyword similarity is assumed to be 1. Yet, we may derive a lowerbound of keyword similarity for $o^{t'}$ and any Q_{next} query in c:

Lemma 2 *Consider an object $o^{t'}$ and a query q in c. If $q \in Q_{next}$, then $SimT(o^{t'}.\psi, q.\psi) > \frac{q.score(k,t)-q.\alpha \cdot UB_{SimS(o^{t'}.c,c)}}{1-q.\alpha}$.*

PROOF. By the definition of Q_{next}, $SimST(o^{t'}, q) > q.score(k,t)$. Because $UB_{SimS(o^{t'}.c,c)}$ is the upperbound of the spatial similarity between the two cells, by replacing $SimS$ in Equation 1, we have $q.\alpha \cdot UB_{SimS(o^{t'}.c,c)} + (1 - q.\alpha) \cdot SimT(o^{t'}.\psi, q.\psi) > q.score(k,t)$, hence deducing the lemma. $\square$

Let $\tau = \min\{ \frac{q.score(k,t)-q.\alpha \cdot UB_{SimS(o^{t'}.c,c)}}{1-q.\alpha} \mid q \in c \}$, i.e., the minimum value among all the queries in c. It is clear that if $SimT(o^{t'}.\psi, q.\psi)$ satisfies the condition in Lemma 2, then it must be greater than τ. Thus, we may use τ as a threshold, and the task of finding the Q_{next} queries in c becomes a *set similarity search* problem [3, 28].

Problem 2 (Set Similarity Search for Q_{next} Queries) *Given an object $o^{t'}$ and a cell c, find the Q_{next} queries in c such that $SimT(o^{t'}.\psi, q.\psi) > \tau$.*

An immediate solution to this problem is to utilize an inverted index, mapping each keyword to a list of queries having this keyword in c, so the queries that share at least one keyword with the object can be identified as candidates to be verified. SKYPE leverages the prefix filtering technique [3, 28] that has been widely adopted for set similarity search. It sorts the keywords of the object by decreasing order of weight, and then only the first few keywords (called prefix, whose length is computed based on the weights of keywords and τ) need to be looked up in the inverted index. This is because if there is no match in the prefix, the other keywords do not have adequate weights to make $SimT(o^{t'}.\psi, q.\psi)$ greater than τ. As we have explained in Section 1, the main drawback of adapting SKYPE to our problem is that the inverted index is built on *single keywords*. Hence the access to the inverted index causes considerable overhead as well as numerous candidates when there are only frequent keywords in $o^{t'}$ (e.g., car and restaurant), which is a common scenario in our problem setting. To address this issue, we propose a signature scheme to efficiently answer the set similarity search in our problem.

Our key idea is to build an inverted index on not only single keywords but also *keyword combinations*. E.g., given a query with two keywords A and B, we index A, B, and AB. We call such keyword combination (including single keywords) *l-signatures*, where l denotes the number of keywords in a signature. In order to control the size of a signature (as the enumeration of the l-signatures of a query is exponential in l), we set a maximum limit of l, denoted by l_{max}. Given a query in a cell c, all the query's $1, \ldots, l_{max}$-signatures are enumerated and indexed; while for $o^{t'}$, we *select* a set of l-signatures and look up in the inverted index. The queries in c having at least one of these l-signatures are regarded as candidates, which will be verified whether they are indeed Q_{next} queries.

Example 3 *Consider three queries q_1, q_2, and q_3 in Figure 3. q_1 has keywords A, B, and C. q_2 has keywords C and D. q_3 has a keyword D. Suppose $l_{max} = 2$. Then the l-signatures of q_1 are*

$\{A, B, C, AB, AC, BC\}$. *The l-signatures of q_2 are $\{C, D, CD\}$. The l-signature of q_3 is $\{D\}$. The corresponding inverted index is shown in Figure 3. Suppose the selected set of l-signatures of $o^{t'}$ is $\{A, CD\}$. Then q_1 and q_2 are the candidates.*

The advantage of using l-signatures results from the observation that an l-signature ($l > 1$) is usually much less frequent than its constituent keywords. E.g., the combination of card and restaurant is much less frequent than either keyword. By looking up the postings list of the combination, the cost of index access and the candidate number can be remarkably reduced.

The signature selection of $o^{t'}$ is essential to the correctness of the algorithm and the query processing performance, since we need to guarantee that no queries satisfying $SimT(o^{t'}.\psi, q.\psi) > \tau$ will be missed and at the same time achieve a small candidate set. For this reason, we propose the notion of object variants.

Definition 4 *Consider $o.\psi$, the keywords of an object o. A variant of o, denoted by v, is a subset of $o.\psi$ such that $SimT(o.\psi, v) > \tau$. The keyword weights of v are normalized to unit length.*

E.g., consider $o.\psi = \{A, B, C\}$, and the weights are 0.332, 0.5, and 0.8, respectively. Suppose $\tau = 0.9$. $\{B, C\}$ is a variant because $SimT(o.\psi, \{B, C\}) = 0.94 > \tau$ (note the normalization of $\{B, C\}$). $\{A, C\}$ is not a variant because $SimT(o.\psi, \{A, C\}) = 0.87 < \tau$. Let $V(o.\psi)$ denote the set of all variants of o. Then the following lemma holds.

Lemma 3 *If $SimT(o.\psi, q.\psi) > \tau$, then $\exists v \in V(o.\psi)$ such that $v \subseteq q.\psi$.*

PROOF. We prove by contradiction. Assume that $\nexists v \in V(o.\psi)$ such that $v \subseteq q.\psi$. Let $r_1 = o.\psi \cap q.\psi$, and the keyword weights are the same as those in $o.\psi$. Let $r_2 = r_1$, and the keyword weights are normalized to unit length. Then $SimT(o.\psi, r_2) \geq SimT(o.\psi, r_1) = SimT(o.\psi, q.\psi) > \tau$. r_2 is a variant of o. Because $r_2 = r_1$, it is a subset of $q.\psi$. This contradicts the assumption. $\square$

The lemma means that the keywords of a Q_{next} query must be a superset of at least one variant of $o^{t'}$. We say an l-signature s *covers* a variant v, iff. the keywords of s are a subset of v; e.g., AB covers ABC but AD does not (we abuse the notations for keyword combinations and sets here). Let $S_{all}(\cdot)$ be the set of all $1, \ldots, l_{max}$-signatures of an object, a query, or a variant. Let $S_{sel}(\cdot)$ be a subset of $S_{all}(\cdot)$. We say that $S_{sel}(o.\psi)$ is a *valid* set of l-signatures of o, iff. $\forall v \in V(o.\psi)$, v is covered by at least one signature $s \in S_{sel}(o.\psi)$. E.g., given a $V(o.\psi) = \{AB, BC, CDE\}$. $\{AB, C\}$ is a valid set of l-signatures, since AB is covered by AB, and BC and CDE are covered by C. Then we have the following relationship between the l-signatures of a query and an object.

Lemma 4 *Consider an object o and $S_{sel}(o.\psi)$, a valid set of l-signatures of o. Given a query q, if $SimT(o.\psi, q.\psi) > \tau$, then $S_{all}(q.\psi) \cap S_{sel}(o.\psi) \neq \emptyset$.*

PROOF. By Lemma 3, if $SimT(o.\psi, q.\psi) > \tau$, $\exists v \in V(o.\psi)$ such that $v \subseteq q.\psi$. Let v^* be one of such v. Because $\forall v \in V(o.\psi)$, v is covered by at least one $s \in S_{sel}(o.\psi)$, $\exists s \in S_{sel}(o.\psi)$, such that v^* is covered by s. Therefore, $s \in S_{all}(v^*)$, and thus $S_{all}(v^*) \cap S_{sel}(o.\psi) \neq \emptyset$. Because $v^* \subseteq q.\psi$, $S_{all}(v^*) \subseteq S_{all}(q.\psi)$. Therefore, $S_{all}(q.\psi) \cap S_{sel}(o.\psi) \supseteq S_{all}(v^*) \cap S_{sel}(o) \neq \emptyset$. $\square$

By this lemma, we can select a valid set of l-signatures for $o^{t'}$. It is guaranteed that any Q_{next} query *must share* at least one

of these selected signatures. For the sake of efficiency, we need to minimize the cost of the selected signatures. The lengths of postings lists in the inverted index are used here to measure the cost. Then the l-signature selection problem is defined below.

Problem 3 (l-signature Selection) *Let $|L_s|$ denote the length of the postings list of a signature s in the inverted index. Given an object o and a threshold τ, the l-signature selection problem is to select $S_{sel}(o.\psi)$, a valid set of l-signatures for o, such that $\sum_{s \in S_{sel}(o.\psi)} |L_s|$ is minimized.*

It is easy to see that the l-signature selection problem is exactly a minimum weighted set cover problem. Hence it is NP-hard and we can solve it by a greedy algorithm with an approximation ratio of $O(ln|V(o.\psi)|)$ [32]. Given variants $V(o.\psi)$ and all the l-signatures of o, the greedy algorithm picks signatures in the order of decreasing benefit-to-cost ratio. The benefit of an l-signature is the number of variants that it covers in the uncovered part of $V(o.\psi)$.

The complexity of the greedy algorithm is $O(|S_{all}(o.\psi)| \cdot |V(o.\psi)|) = O(\binom{|o.\psi|}{l_{max}} \cdot (2^{|o.\psi|} - 1))$. $|o.\psi|$ is usually small in real applications, and we can choose a small l_{max} to control the enumeration cost of $S_{all}(o.\psi)$.

Enumerating variants. Prior to running the greedy algorithm, a technical challenge is how to compute $V(o.\psi)$. One may notice that we only need to cover the minimal variants (i.e., the variant such that we cannot remove any keyword from it and still make a variant) in $V(o.\psi)$ instead of all. However, the minimality check is costly. We take a compromise to quickly find a subset of $V(o.\psi)$ and still guarantee the selected l-signatures cover $V(o.\psi)$: First, we sort the keywords of $o.\psi$ by decreasing order of weights, and pick keywords one by one until we get a variant. It can be seen that the number of picked keywords is the minimum number of keywords of a variant. So we only enumerate subsets of $o.\psi$ whose sizes are no smaller than this number. By enumerating subsets of $o.\psi$ in increasing order of size, we can stop when reaching a size such that all the subsets of this size are variants. Let $V'(o.\psi)$ denote the set of variants we have enumerated. We replace $V(o.\psi)$ with $V'(o.\psi)$ and run the greedy algorithm. The resulting selected l-signatures are guaranteed to cover $V(o.\psi)$, because any variant in $V(o.\psi)$ is a superset of at least one variant in $V'(o.\psi)$. E.g., suppose there are 5 keywords in $o.\psi$, and the minimum number of keywords in a variant is 2. So we enumerate subsets of $o.\psi$ whose sizes are 2, 3, ... Suppose all the subsets of size 3 are variants. We stop at 3 and $V'(o.\psi)$ includes only variants of size 2 and 3.

By putting it all together, the pseudo-code of the processing of each cell is provided in Algorithm 2. It first computes the τ threshold of the cell (Line 1). The V' variants of $o^{t'}$ are generated (Line 2). We enumerate the l-signatures of $o^{t'}$ (Line 3) and select the ones that cover V' (Line 4). For each selected signature, we access the corresponding postings list of the inverted index of this cell (Line 7). Each query in the list is a candidate and inserted into Q_{next} if the score between $o^{t'}$ and the query exceeds the k-th result at t (Line 8). We put the pseudo-code of the greedy signature selection algorithm in the full version of this paper [11] and omit it here.

One signature selection for all. When invoking Algorithm 2 in Algorithm 1, a limitation is that the signature selection algorithm is run for each cell connected by a CC link and thus causes redundancy. We observe that the variant set $V(o.\psi)$ only depends on the object and the keyword similarity threshold τ. Given two cells c_1 and c_2 with thresholds τ_1 and τ_2, respectively,

---

Algorithm 2: GetQueryIn$(c, o^{t'})$

---

1 $\tau \leftarrow \min\{ \frac{q.score(k,t)-q.\alpha \cdot UB_{SimS(o^{t'}.c,c)}}{1-q.\alpha} \mid q \in c \}$;

2 $V' \leftarrow$ GenerateVariants$(o^{t'}, \tau)$;

3 $S_{all} \leftarrow$ all l-signatures of $o^{t'}$;

4 $S_{sel} \leftarrow$ SelectSignatures(V', S_{all});

5 $Q_{next} = \emptyset$;

6 **foreach** $s \in S_{sel}$ **do**

7   **foreach** $q \in c.L_s$ **do**

8   **if** $SimST(o^{t'}, q) > q.score(k,t)$ **then**

             $Q_{next} \leftarrow Q_{next} \cup \{ q \}$;

9 **return** Q_{next}

Algorithm 3: SelectSignatures(V, S_{all})

---

1 $S_{sel} \leftarrow \emptyset$;

2 **foreach** $s \in S_{all}$ **do**

3 $cost_s \leftarrow |L_s|$;

4 **while** $V \neq \emptyset$ **do**

5 $cover_s \leftarrow$ the set of variants in V covered by s;

6 $benefit_s \leftarrow |cover_s|$;

7 $S_{sel} \leftarrow S_{sel} \cup \{s \mid s \in S_{all} \wedge s$ has maximum $benefit_s/cost_s \}$;

8 $S_{all} \leftarrow S_{all}\backslash\{ s \}, V \leftarrow V\backslash\{ cover_s \}$;

9 **return** S_{sel}

---

it can be seen that the corresponding variant sets $V_1(o.\psi) \subseteq V_2(o.\psi)$ if $\tau_1 \geq \tau_2$. This suggests the l-signatures selected using $V_2(o.\psi)$ can cover $V_1(o.\psi)$ as well. Hence, we may take the minimum τ across all the cells identified by CC links, and run the greedy algorithm only *once* to obtain the l-signatures for all these cells: First, we generate the variant set $V'(o.\psi)$ by the minimum τ. While running the greedy algorithm, because $V_1(o.\psi) \subseteq V_2(o.\psi)$, $V_1(o.\psi)$ is always covered prior to $V_2(o.\psi)$. So we monitor for each cell the time when the corresponding variant set has just been covered, and record the l-signatures selected so far as the l-signatures for this cell.

We briefly discuss the differences of l-signatures from existing studies, since keyword (token) combinations were also used for approximate set containment search [1] and local similarity search [23]: (1) [1] enumerates keyword combinations frequency, while we use number of keywords to control the enumeration. (2) [1] covers minimal variants, while we choose a compromise between minimal and all. (3) [23] is based on prefix filtering and partitioning. It is not applicable for low thresholds since the prefix length may exceed the size of an object. Our method does not have such limitation. (4) We propose the one-selection-for-all technique tailored to our problem setting.

6. TOP-K REFILLER

After identifying Q_{prev} and Q_{next} by the AQF, we update the top-k results for the queries in the two sets. For the queries in Q_{next}, the update is straightforward because we only need to insert $o^{t'}$ into the query's top-k and delete the previous k-th object from it. For the queries in Q_{prev}, since $o^{t'}$ may move out of the top-k, we need to refill the list with another object. A sequential scan of all the objects in O is prohibitively expensive for online query processing. One option is to borrow the top-k refilling method in kmax [31]. It maintains a list of objects in

a buffer that can be quickly filled as top-k. As discussed in Section 1, keeping objects in a buffer is inefficient in our problem setting because objects are dynamic, incurring frequent buffer maintenance and hence considerable overhead. Next we propose two methods to deal with top-k refilling.

6.1 The Cell Bound (CB) Method

We explore in the direction of exploiting the score bounds of cells. It has the following advantages: (1) the number of cells is much smaller than the number of objects (compared to the buffering strategy in kmax); (2) the bound of spatial similarity from a cell to a query is static; and (3) despite objects being dynamic, the bounds of keyword similarities from the objects in a cell to a query do not change frequently.

The bound of spatial similarity from a cell to a query is easy to derive. To bound keyword similarities, additional information needs to be stored in the grid index. For each cell c, we collect the distinct keywords that appears in the objects in c at t, denoted by $c^t.\psi$. Then we have two weights for each keyword, the maximum and the minimum weights of w among all the objects in c that have w at t. For simplicity, we use $c^t.\psi_{max}$ and $c^t.\psi_{min}$ to distinguish the weight we want to use, though they refer to the same set of keywords $c^t.\psi$. When an update of an object occurs, such information is updated accordingly.

Example 4 *In Figure 3, there are two objects, o_1 and o_2, in c_1. Suppose there are three distinct keywords in o_1 and o_2: A, B, and C. Then $c_1^t.\psi = \{ A, B, C \}$. The weights in $c_1^t.\psi_{max}$ are $\{ 0.2, 0.4, 0.1 \}$. The weights in $c_1^t.\psi_{min}$ are $\{ 0.2, 0.3, 0.1 \}$.*

With $c^t.\psi_{max}$ and $c^t.\psi_{min}$, the upper and lower bounds of keyword similarity in c are respectively bounded by $SimT(c^t.\psi_{max}, q.\psi)$ and $SimT(c^t.\psi_{min}, q.\psi)$, as computed by Equation 3. Then we can bound the scores of the objects in a cell.

$$maxscore(c^t, q) = q.\alpha \cdot UB_{SimS(c,q)} + (1-q.\alpha) \cdot UB_{SimT(c^t,q)}. \tag{5}$$

$$minscore(c^t, q) = q.\alpha \cdot LB_{SimS(c,q)} + (1-q.\alpha) \cdot LB_{SimT(c^t,q)}. \tag{6}$$

UB_{SimS} and LB_{SimS} are upper and lower bounds of spatial similarity, respectively: $UB_{SimS(c,q)} = 1 - \frac{minD(c,q.\rho)}{maxDist}$, $LB_{SimS(c,q)} = 1 - \frac{maxD(c,q.\rho)}{maxDist}$, where $minD(c, q.\rho)$ and $maxD(c, q.\rho)$ denote the minimum and the maximum distances from a cell to the query location, respectively. UB_{SimT} and LB_{SimT} are upper and lower bounds of keyword similarity, respectively: $UB_{SimT(c^t,q)} = SimT(c^t.\psi_{max}, q.\psi)$ and $LB_{SimT(c^t,q)} = \min \{ wt(c^t.\psi_{min}, w) \cdot wt(q.\psi, w) \mid w \in c^t.\psi_{min} \cap q.\psi \}$.

We design a top-k refilling algorithm by leveraging the cell bounds (CB). The pseudo-code is shown in Algorithm 4. For the queries in Q_{next}, we only insert $o^{t'}$ into the queries' top-k (Line 2). Here, we assume that $q.obj(1..k, t)$ is a priority queue of size k, whose elements are ranked by the score. The $push()$ function inserts an object to the priority queue and only keeps the top-k ones. For the queries in $Q_{prev}\backslash Q_{next}$ (the set difference is because the overlapping queries have been processed in Line 2), we remove o from the top-k (Line 4) [4] and look for the k-th result at t'. The $maxscore(c^{t'}, q)$ and

---

[4]A special case is $SimST(o^{t'}, q) = q.score(k,t)$. o is exactly the k-th result of q at t' and no further action is required. We omit this case in the pseudo-code for conciseness.

Algorithm 4: TopKRefiller-CB($o^t, o^{t'}, OutQ, InQ$)

    Input   :the stati of o at t and t', Q_{prev}, Q_{next}
    Output:top-k results for $q \in Q_{prev} \cup Q_{next}$

1 **foreach** $q \in Q_{next}$ **do**
2 $\lfloor$ $q.obj(1 \mathinner{.\,.} k, t') \leftarrow q.obj(1 \mathinner{.\,.} k, t).push(o)$;

3 **foreach** $q \in Q_{prev} \backslash Q_{next}$ **do**
4 $q.obj(1 \mathinner{.\,.} k, t') \leftarrow q.obj(1 \mathinner{.\,.} k, t).pop(o)$;
5   **foreach** $c \in C$ **do**
6 $\lfloor$ Compute $maxscore(c^{t'}, q)$ and $minscore(c^{t'}, q)$;

7   Sort C by descending $maxscore(c^{t'}, q)$;
8 $score_k \leftarrow -\infty$;
9   **foreach** $c \in C$ **do**
10   **if** $maxscore(c^{t'}, q) \leq score_k$ **then break**;
11   **if** $minscore(c^{t'}, q) > q.score(k, t)$ **then**
           continue;
12   **foreach** $o' \in c$ such that $o' \notin q.obj(1 \mathinner{.\,.} k, t')$ **do**
13   **if** $SimST(o'^{t}, q) > score_k$ **then**
14 $\lfloor$ $o_k \leftarrow o'$, $score_k \leftarrow SimST(o'^{t}, q)$;

15 $q.obj(1 \mathinner{.\,.} k, t').push(o_k)$;

16 **return** $\{ q.obj(1 \mathinner{.\,.} k, t') \mid q \in Q_{prev} \cup Q_{next} \}$

$minscore(c^{t'}, q)$ values of all the cells are updated (Line 6), and then the cells are sorted by descending order of $maxscore$ (Line 7). We keep a temporary object (o_k) for the k-th result at t' and its score ($score_k$) when iterating through the cells in the grid. The objects in the cells are scanned and the temporary object is updated (Lines $12-14$). Finally, the temporary object with the highest score is inserted into the top-k of q (Line 15). There are two optimizations: First, if $maxscore$ is no better than the temporary object's score (Line 10), the iteration is early terminated because unseen cells are even lower in $maxscore$. Second, if $minscore$ is higher than the k-th result at t (Line 11), we scan skip this cell, as it can be shown that $q.obj(k, t')$ never has a higher score than $q.obj(k, t)$ for any q in $Q_{prev} \backslash Q_{next}$ (because otherwise q must be in Q_{next}).

Example 5 *We consider the processing of a query in $Q_{prev} \backslash Q_{next}$. Figure 4 (top left corner) shows the cells sorted by descending maxscore. Suppose $q.score(k, t) = 0.7$. The first cell c_5 is skipped since its minscore is greater than $q.score(k, t)$. Suppose the temporary object's score becomes 0.5 when we reach c_8. We can stop here because c_8's maxscore does not exceed 0.5. Then the temporary object is refilled to the top-k of q at t'.*

6.2   The Cell List (CL) Method

The main drawback of the CB method is that it has to scan all the cells in the grid for each query; however, it can be seen that when o^t changes to $o^{t'}$, at most one object in the top-k of q changes. Moreover, we have the following observation.

Observation 1 *Consider an old status o^t and a new status $o^{t'}$ of an object. For any $q \in Q_{prev}$, $q.o(1 \mathinner{.\,.} k, t') \backslash q.obj(1 \mathinner{.\,.} k, t) = \{ o \}$ or $\{ q.obj(k + 1, t) \}$.*

$q.obj(k + 1, t)$ is the $(k + 1)$-th result of q at time t. The observation indicates that we only need to compare $o^{t'}$ and $q.obj(k + 1, t)$, and pick the one with the higher score as the new top-k result of q. Based on this observation, we devise an improved method for top-k refilling by maintaining a cell list

for each q such that the $(k + 1)$-th object of q at t is guaranteed to reside in one of these cells [5]. Like the CB method, our idea is also to exploit the score bounds, yet in a more effective way. We have the following two lemmata.

Lemma 5 *If $minscore(c^t, q) > q.score(k, t)$, then $q.obj(k + 1, t) \notin c^t$.*

PROOF. Because $minscore(c^t, q) > q.score(k, t)$, we have $\forall o_i \in c^t$, $SimST(o_i^t, q) > q.score(k, t) \geq q.score(k + 1, t)$. Therefore, $q.obj(k + 1, t) \notin c^t$. $\square$

Let $maxminscore_{<k}(C^t, q) = \max\{ minscore(c^t, q) \mid c \in C \land maxscore(c^t, q) < q.score(k, t) \}$; i.e., we collect the cells whose $maxscore$ values are less than $q.score(k, t)$, and pick the the maximum of the $minscore$ values of these cells.

Lemma 6 *If $maxscore(c^t, q) < maxminscore_{<k}(C^t, q)$, then $q.obj(k + 1, t) \notin c^t$.*

PROOF. We prove by contradiction: assume $q.obj(k+1, t) \in c^t$. Let c^* be the cell that yields $maxminscore_{<k}(C^t, q)$; i.e., $maxminscore_{<k}(C^t, q) = minscore(c^{*^t}, q)$. $maxscore(c^t, q) < maxminscore_{<k}(C^t, q)$, so $maxscore(c^t, q) < minscore(c^{*^t}, q)$. Because $maxscore$ and $minscore$ are scores' upper and lower bounds, respectively, and $q.obj(k+1, t) \in c^t$, we have $\exists o' \in c^{*^t}$, s.t. $SimST(o'^{t}, q) > q.score(k + 1, t)$. By the definition of $maxminscore_{<k}$, $maxscore(c^{*^t}, q) < q.score(k, t)$. Therefore, $SimST(o'^{t}, q) < q.score(k, t)$, meaning that $SimST(o'^{t}, q) \leq q.score(k+1, t)$. This contradicts $SimST(o'^{t}, q) > q.score(k+1, t)$. $\square$

Intuitively, the two lemmata state that if a cell whose score bounds are too high or too low, then the $(k + 1)$-th result is not in it. So by taking the complement of these cells in the grid, we are guaranteed to include the $(k + 1)$-th result's cell:

Corollary 1 *Consider a set of cells $C' = \{ c \in C \mid minscore(c^t, q) \leq q.score(k, t) \land maxscore(c^t, q) \geq maxminscore_{<k}(C^t, q) \}$. $\exists c \in C'$, s.t., $q.obj(k + 1, t) \in c^t$.*

PROOF. We prove by contradiction. If $\nexists c \in C'$, s.t., $q.obj(k+1, t) \in c^t$, by the definition of C', the cell that has $q.obj(k+1, t)$, denoted by c', satisfies either $minscore(c'^{t}, q) > q.score(k, t)$ or $maxscore(c'^{t}, q) < maxminscore_{<k}(C^t, q)$. This either contradicts Lemma 5 or Lemma 6. $\square$

Although it is sufficient to fetch the $(k+1)$-th result with only C' in Corollary 1, we also store the cells having the top-$1, \ldots, k$ results. This is to simplify the update of the cell list (introduced later). As such, we define the $(k + 1)$-CL of q:

Definition 5 ($(k + 1)$-CL) *The $(k + 1)$-CL of a query q with a timestamp t is defined as $\{ c \in C \mid maxscore(c^t, q) \geq maxminscore_{<k}(C^t, q) \}$.*

Compared to Corollary 1, we remove the condition on $minscore$. In the rest of this section, we mean $(k + 1)$-CL when CL is mentioned, and we use $q.CL$ to denote the CL of q. Since we will use a lazy update strategy, the corresponding timestamp is recorded along with the CL, denoted by $q.CL.t$.

Example 6 *In Figure 4 (time t), cells are represented by intervals. The two ends indicate the $maxscore(c^t, q)$ and $minscore(c^t, q)$ values of a cell. By the definition of $(k+1)$-CL, $q.CL$ with timestamp t is $\{ c_5, c_2, c_1, c_6 \}$.*
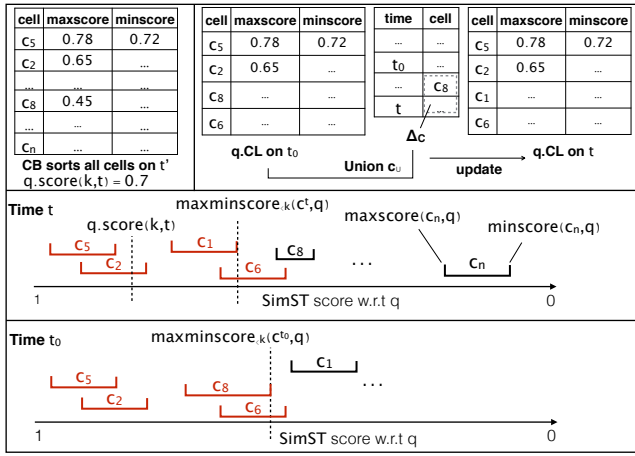
Figure 4: top-k refilling.

Algorithm 5: TopKRefiller-CL($o^t, o^{t'}, OutQ, InQ$)

Input : the stati of o at t and t', Q_{prev}, Q_{next}
Output : top-k results for $q \in Q_{prev} \cup Q_{next}$

1 **for** each $q \in Q_{next}$ **do**
2 $\quad$ $q.obj(1\mathbin{..}k, t') \leftarrow q.obj(1\mathbin{..}k, t).push(o)$;

3 **foreach** $q \in \{Q_{prev} \backslash Q_{next}\}$ **do**
4 $\quad$ $q.obj(1\mathbin{..}k, t') \leftarrow q.obj(1\mathbin{..}k, t).pop(o)$;
5 $\quad$ UpdateCL($q.CL, q, t$);
6 $\quad$ $q.CL_{\leq k} \leftarrow \{ c \mid c \in q.CL \land minscore(c^t, q) \leq q.score(k, t) \}$;
7 $\quad$ Sort $q.CL_{\leq k}$ by descending $maxscore(c^t, q)$;
8 $\quad$ $score_{k+1} \leftarrow -\infty$;
9 $\quad$ **foreach** $c \in q.CL_{\leq k}$ **do**
10 $\quad\quad$ **if** $maxscore(c^t, q) \leq score_{k+1}$ **then break**;
11 $\quad\quad$ **foreach** $o' \in c$ such that $o' \neq q.obj(k, t)$ **do**
12 $\quad\quad\quad$ **if** $q.score(k, t) \geq SimST(o'^t, q) > score_{k+1}$
$\quad\quad\quad\quad$ **then**
13 $\quad\quad\quad\quad$ $o_{k+1} \leftarrow o', score_{k+1} \leftarrow SimST(o'^t, q)$;

14 $\quad$ **if** $SimST(o^{t'}, q) \geq score_{k+1}$ **then**
$\quad\quad$ $q.obj(1\mathbin{..}k, t').push(o)$;
15 $\quad$ **else** $q.obj(1\mathbin{..}k, t').push(o_{k+1})$;

16 **return** $\{ q.obj(1\mathbin{..}k, t') \mid q \in Q_{prev} \cup Q_{next} \}$

We design an improved top-k refilling algorithm utilizing CL. Algorithm 5 gives the pseudo-code. Like in the CB algorithm, for any query in Q_{next}, we only need to push $o^{t'}$ into the query's top-k (Line 2). For any query in $Q_{prev} \backslash Q_{next}$, we first pop o from the top-k (Line 4), and update the query's CL to the state at t (not t', Line 5). Then we use the part of the CL that meet the condition in Corollary 1 (Line 6) to find the $(k+1)$-th result. The cells in this part are sorted (Line 7). Then the algorithm follows a similar procedure (Lines $8 - 13$) with the cell iteration in Algorithm 4. The $(k+1)$-th result at t is kept in a temporary object o_{k+1} along with its score $score_{k+1}$. Finally, we compare o_{k+1} with o and insert the one with the higher score to the top-k results of q (Lines $14 - 15$).

---

[5] One may want to keep only the $(k + 1)$-th object, but this object may change status while q is outside both Q_{prev} and Q_{next}, hence difficult to track.

Example 7 *We consider a query q in $Q_{prev} \backslash Q_{next}$. In Figure 4, the cell list of q with timestamp t includes c_5, c_2, c_1, and c_6. Since c_5 violates the condition in Corollary 1 (due to $minscore(c^t, q)$ exceeding $q.score(k, t)$), we only access c_2, c_1, and c_6 to retrieve the $(k + 1)$-th result at t. Then it is compared with $o^{t'}$. The one with the higher score is refilled to the top-k of q at t'.*

6.2.1 Cell List Update

An important step in Algorithm 5 is the update of CL (Line 5), as we need to guarantee the CL has all the cells satisfying the condition in Definition 5. A naive way of update is to scan all the cells to make a CL from scratch, but this essentially reduces the algorithm to the CB method. An efficient solution is to find the cells to be inserted or deleted from the existing CL. The challenge here is that the CL's timestamp (say, t_0) prior to the update could be very old compared to t. Nonetheless, by Definition 5, the cells in the CL are determined by $maxscore$, $minscore$, and $q.score(k, t)$. $maxscore$ and $minscore$ depend only on the keywords that appear in the cell. Hence, we may keep track of the cells whose $c.\psi_{\max}$ or $c.\psi_{\min}$ change, along with the timestamp at which the change happens. Then the new CL can be computed using the old CL and the set of the cells that change $c.\psi_{\max}$ or $c.\psi_{\min}$ during time interval $(t_0, t]$. Let Δ_C denote this set. We describe our solution to CL update.

The key idea is to compute $maxminscore_{<k}(C^t, q)$, so the inserted/deleted cells can be determined by comparing their $maxscore$ values with $maxminscore_{<k}(C^t, q)$. To this end, we derive the following lemma.

Lemma 7 *Consider $q.CL$, a cell list to be updated, whose timestamp is t_0. Let $C_\cup = q.CL \cup \Delta_C$. If $maxminscore_{<k}(C_\cup^t, q) \geq maxminscore_{<k}(C^{t_0}, q)$, then we have $maxminscore_{<k}(C^t, q) = maxmins(C_\cup^t, q)$.*

PROOF. By the definition of $maxminscore_{<k}$, for any $C' \subseteq C$, $maxminscore_{<k}(C^t, q) \geq maxminscore_{<k}(C'^t, q)$. Because $C_\cup \subseteq C$, $maxminscore_{<k}(C^t, q) \geq maxminscore_{<k}(C_\cup^t, q)$. Therefore, if $maxminscore_{<k}(C^t, q) \neq maxminscore_{<k}(C_\cup^t, q)$, $maxminscore_{<k}(C^t, q) > maxminscore_{<k}(C_\cup^t, q)$. We assume this is true and prove by contradiction. In this case, $\exists c' \in C \backslash C_\cup$, s.t., $maxscore(c'^t, q) > maxminscore_{<k}(C_\cup^t, q)$. Because $maxminscore_{<k}(C_\cup^t, q) \geq maxminscore_{<k}(C^{t_0}, q)$, and any cell in $C \backslash C_\cup$ does not change $maxscore$ or $minscore$ during $(t_0, t]$, we have $maxscore(c'^t, q) = maxscore(c'^{t_0}, q) > maxminscore_{<k}(C^{t_0}, q)$. This means $c' \in q.CL$ and thus $c' \in C_\cup$, which contradicts $c' \in C \backslash C_\cup$. $\square$

The lemma reveals under which condition we can make the new CL using the old CL and Δ_C. First, we update the $maxscore$ and $minscore$ values of the cells in $q.CL$ and Δ_C for timestamp t, and compute $maxminscore_{<k}(C_\cup^t, q)$ using these $maxscore$ and $minscore$ values and $q.score(k, t)$. Then we check if $maxminscore_{<k}(C_\cup^t, q)$ is no less than that of the old CL, i.e., $maxminscore_{<k}(C^{t_0}, q)$. If so, $maxminscore_{<k}(C^t, q)$ is exactly $maxminscore_{<k}(C_\cup^t, q)$. All the cells to be inserted or deleted must belong to $C_\cup$. This is because the cells having $maxscore(c^t, q) \geq maxminscore_{<k}(C^{t_0}, q)$ are a superset of those having $maxscore(c^t, q) \geq maxminscore_{<k}(C_\cup^t, q) = maxminscore_{<k}(C^t, q)$, and the former are stored in $C_\cup$. Thus, the new CL can be made using only $C_\cup$. Otherwise (i.e., when $maxminscore_{<k}(C_\cup^t, q) < maxminscore_{<k}(C^{t_0}, q)$), we scan all the cells of the grid to make the new CL.

The pseudo-code of the above process is captured by Algorithm 6. Lines $1 - 4$ update $maxscore$ and $minscore$ in $C_\cup$

Algorithm 6: UpdateCL($q.CL, q, t$)

---

1 $\Delta_C \leftarrow$ the cells that change $c.\psi_{\max}$ or $c.\psi_{\min}$ in $(q.CL.t, t]$;
2 **foreach** $c \in \Delta_C$ **do**
3 $\quad$ Compute $maxscore(c^t, q)$ and $minscore(c^t, q)$;
4 $C_\cup \leftarrow q.CL \cup \Delta_C, \gamma \leftarrow maxminscore_{<k}(C_\cup^t, q)$;
5 **if** $\gamma \geq maxminscore_{<k}(C^{q.CL.t}, q)$ **then**
6 $\quad$ **foreach** $c \in q.CL$ **do**
7 $\quad\quad$ **if** $maxscore(c^t, q) < \gamma$ **then** $q.CL.delete(c)$;
8 $\quad$ **foreach** $c \in \Delta_C \backslash q.CL$ **do**
9 $\quad\quad$ **if** $maxscore(c^t, q) \geq \gamma$ **then** $q.CL.insert(c)$;
10 **else**
11 $\quad$ $q.CL.clear(), \gamma \leftarrow maxminscore_{<k}(C^t, q)$;
12 $\quad$ **foreach** $c \in C$ **do**
13 $\quad\quad$ **if** $maxscore(c^t, q) \geq \gamma$ **then** $q.CL.insert(c)$;
14 $q.CL.t \leftarrow t$;

---

and compute $maxminscore_{<k}(C_\cup^t, q)$. Lines 5 – 9 make the new CL using the cells in $C_\cup$ only. Lines 11 – 13 make the new CL by scanning all the cells in the grid. The timestamp of the CL is updated to t eventually (Line 14).

Example 8 *As shown in Figure 4, we assume that the old $q.CL$ with timestamp t_0 is $\{ c_5, c_2, c_8, c_6 \}$. $maxminscore_{<k}(C^{t_0}, q)$ is given by $minscore(c_6^{t_0}, q)$. Suppose $\Delta_C = \{ c_1, c_8 \}$. So $C_\cup = \{ c_5, c_2, c_8, c_6, c_1 \}$. $maxminscore_{<k}(C_\cup^t, q)$ is given by $minscore(c_1^t, q)$. Since it is greater than $maxminscore_{<k}(C^{t_0}, q)$, $maxminscore_{<k}(C^t, q) = maxminscore_{<k}(C_\cup^t, q)$. So only the cells in $C_\cup$ are considered to compute the new CL. We compare their $maxscore(c^t, q)$ values with $maxminscore_{<k}(C^t, q)$. c_8 is deleted and c_1 is inserted. The new CL is $\{ c_5, c_2, c_1, c_6 \}$.*

7. BATCH PROCESSING

We extend our method to batch processing. Given the top-k results of each query of Q at t, our task is to compute the top-k results at t'. Let $O^{t,t'}$ denote the set of objects that change their stati during time interval $(t, t']$.

7.1 Affected Query Finder for Batch

We first modify the definitions of Q_{prev} and Q_{next}: (1) Q_{prev}: the multiset of queries such that $\exists o \in O^{t,t'}$, o is a top-k result of q at t, and (2) Q_{next}: the multiset of queries such that $\exists o \in O^{t,t'}$, $SimST(o^{t'}, q) > q.score(k, t)$. The occurrence of a query q in Q_{prev}/Q_{next} is the number of objects in $\in O^{t,t'}$ satisfying the above condition. The top-k results are not affected if a query is in neither Q_{prev} nor Q_{next}.

The computation of Q_{prev} is straightforward. We scan the objects in $O^{t,t'}$ and fetch Q_{prev} queries by the object-query map. To find Q_{next} queries, we scan the objects in $O^{t,t'}$ and use CC links to identify the cells that need look-up. An important observation is that objects may share keywords and computation can be shared when invoking Algorithm 2. Consider a cell c identified by CC links from multiple objects in $O^{t,t'}$. Each object o yields its own keyword similarity threshold. We generate the l-signature set S_{all} and the variant set V' for each object. Since these l-signature sets and variant sets may overlap, we take them together as the input of the signature selection algorithm. While running the signature selection, we monitor the time when the variant set of an object has been covered, and record the

l-signatures selected so far. This resembles the one-signature-selection-for-all technique described in Section 5.2. Besides, the two techniques can work together. As such, only one run of signature selection outputs the l-signatures for all the $O^{t,t'}$ objects and all the cells identified by CC links.

7.2 Top-k Refiller for Batch

Since Q_{prev} and Q_{next} are multisets and may overlap, we need to count for each query q its occurrences in Q_{prev} and Q_{next} to determine what objects need to be refilled to top-k. Let Δ_q denote the difference of q's occurrences in Q_{prev} and Q_{next}. If $\Delta_q \leq 0$, the processing is similar to Q_{prev} queries in the case of single object update: we pop out the Q_{prev} objects (i.e., the $O^{t,t'}$ objects which are q's top-k results at t), and push into q's top-k the Q_{next} objects (i.e., the $O^{t,t'}$ objects such that $SimST(o^{t'}, q) > q.score(k, t)$). If $\Delta_q > 0$, the processing is similar to Q_{next} queries in the case of single object update: (1) When the the CB method is chosen, the $maxscore$ and $minscore$ bounds of each cell are updated first, and we scan all the cells to find the top-Δ_q objects to be refilled. Then the Q_{prev} objects are popped out, the Q_{next} objects and the top-Δ_q objects are pushed into q's top-k. (2) When the CL method is chosen, the Q_{prev} objects are popped out first. The CL of q has to include the $2k$-th object at t, because Δ_q can be up to k. To this end, $maxminscore_{<k}(C^t, q)$ is replaced by $kminscore_{<k}(C^t, q)$, which is k-th largest $minscore(c^t, q)$ of the cells in C such that $maxscore(c^t, q) < q.score(k, t)$. Definition 5 is replaced by $(2k)$-CL. The related lemmata, corollary, and algorithms are modified accordingly. Then we use the $(2k)$-CL to fetch the $(k + 1)$-th to $2k$-th objects of q at t. They are compared with the Q_{prev} objects and the better ones are pushed into q's top-k, along with the Q_{next} objects.

8. INITIALIZATION AND QUERY UPDATES

Initialization. We first create a grid index on the objects. The initialization of top-k results for each query is a snapshot spatial-keyword search. We do not choose existing methods (surveyed in Section 2) as they either target different problem settings or use different indexes. Our method maintains a list of temporary top-k results (initialized as empty). By starting from the query's cell, it iterates cells by a breadth-first-search. The objects in each cell are checked whether they outperform any temporary top-k result. This is repeated until we reach the cells that are too far to make a score better than the current k-th one. After obtaining top-k results, the query is indexed in the grid. We create an object-query map for each object and a cell list (for the CL method) for each query. Then we create the (l-signature) inverted index and CC links for each cell, and compute the keyword similarity threshold τ.

Query insertion. Like initialization, we also use snapshot spatial-keyword search to retrieve the top-k results of the inserted query. The difference is that the insertion is an online operation, so we need to optimize the processing speed. For the above snapshot spatial-keyword search method, a good set of initial top-k results help prune unpromising cells and objects. An observation is that the top-k results of two similar queries tend to resemble. For this reason, we first find the nearest neighbor query of the inserted query, and use the nearest neighbor's top-k results as an initial set of top-k for the inserted query. Then we run the snapshot spatial-keyword search. Finally, the

Table 3: Datasets statistics.

Datasets	YELP	TWITTER	SYN
Data size	1.2M	4.2M	12M
Default # of objects	220K	500K	1M
Default # of queries	192K	250K	1M
Total # of keywords	819K	3.5M	6M
# of kw. per object/query	5.9	4.5	3

(a) Varying n. (b) Varying $l_{\max}$.

Figure 5: Parameter Tuning.

cell list of the inserted query is created and corresponding data structures (grid, object-query map, etc.) are updated.

Query deletion. The query and its cell list are deleted. The grid, the object-query map, and the inverted index of the query's cell are updated. If the query is the one that determines the CC links or τ of its cell, the CC links or τ needs recomputation.

9. EXPERIMENTS

9.1 Setting

Datasets. We used two real datasets and one synthetic dataset. Table 3 shows the statistics.

- **YELP** is a public dataset with over 192K businesses' information and 1.2M reviews from the Yelp business directory service [30]. We extracted the businesses' locations and descriptions as queries. Users were regarded as objects. The keywords of an object were extracted from the reviews. Since locations are not covered by the reviews, we paired each review with the locations in a real taxi trip data [19].
- **TWITTER** is a dataset with 4.2M geotagged tweets from 1.2M users [22]. We randomly selected a subset of users as queries, using their locations and 1 to 5 keywords randomly chosen from their first tweets. Then we randomly selected a subset of users from the rest of the dataset as objects, locations and keywords extracted in the same way as queries.
- **SYN** is a synthetic data containing 12M spatial keyword tuples. We used a dataset of moving points generated by the BerlinMOD benchmark [12] for locations, and randomly generated keywords by a Zipfian distribution.

Algorithms. We consider the following methods for affected query finder (AQF): (1) CIQ, the inverted file method in [5]; (2) SKYPE, the prefix filtering method in [24]; and (3) CCLS, our method that utilizes **CC** links and *l*-sigs. We consider the following methods for top-k refiller (TR): (1) kmax, the method with $kmax$ buffer in [31]; (2) CB, our cell bound method; and (3) CL, our cell list method. We use "A-B" to denote the combination of methods A and B, e.g., CIQ-kmax.

Environments and measures. The experiments were run on a MacBook Pro with a 2.2GHz Intel Core i7 CPU and 32GB memory. All the algorithms were implemented in C++ and in an main memory fashion. We randomly monitored 10,000 updates of dynamic objects and reported the average processing time per object update. The default k is 20. The default $q.\alpha$ of each query is a random value in $[0, 1]$ by uniform distribution.

9.2 Parameter Tuning

There are two parameters in our methods: the grid size and $l_{\max}$, the maximum number of keywords in an l-signature. We use an $n \times n$ grid here. Figure 5a shows how the running times of our methods vary with n. The grid size mainly affects CC links and the pruning power in TR. The best choice of n is around


(a) Average processing time. (b) Memory consumption.

Figure 6: Overall performance.

(a) # of queries accessed. (b) Average processing time.

Figure 7: Affected query finder.

(a) # of objects accessed. (b) Average processing time.

Figure 8: Top-k refiller.

(a) YELP. (b) TWITTER.

Figure 9: Varying k.

16 to 20. Too few cells reduce the pruning power, while too many cells result in more cell access. Figure 5b shows how $l_{\max}$ affects the overall performance. When $l_{\max} = 1$, our method is the same as indexing single keywords. It can be seen that using keyword combinations is better than single keywords. The best choice is $l_{\max} = 3$. When we use too large $l_{\max}$, the performance drops due to more enumeration cost. We choose $n = 20$ and $l_{\max} = 3$ in the rest of the experiments.
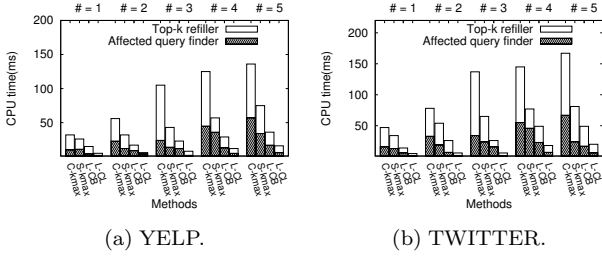
(a) YELP.

(b) TWITTER.

Figure 10: Varying number of keywords in queries.

(a) TWITTER.

(b) SYN.

Figure 14: Varying number of objects $|O|$.

(a) YELP.

(b) TWITTER.

Figure 11: Varying α.

(a) TWITTER, sum = 2M.

(b) SYN, sum = 5M.

Figure 15: Varying ratio of $|Q|$ to $|O|$.

(a) SYN.

(b) SYN.

Figure 12: Effect on object change.

(a) YELP.

(b) TWITTER.

Figure 16: Batch processing.

(a) TWITTER.

(b) SYN.

Figure 13: Varying number of queries $|Q|$.

9.3 Query Processing Performance

We compare our methods (CCLS-CB and CCLS-CL) with alternative solutions and plot the average query processing times on the three datasets in Figure 6a. CIQ-kmax is the slowest. By using prefix filtering, SKYPE-kmax performs better than CIQ-kmax. Both CCLS-CB and CCLS-CL outperforms the two alternative solutions. The speed-up of CCLS-CB over SKYPE-kmax ranges from 1.8 to 2.3 times. CCLS-CL further improves the speed-up to a range of 5.3 to 5.9 times, hence reducing the average query processing time to a range of 8.1ms to 14.5ms. As for memory consumption, we plot the results in Figure 6b. CIQ-kmax and SKYPE-kmax include quadtree, inverted index on single keywords, and object buffer. Our methods include grid, inverted index on l-signatures, and cell list (CL only). The memory consumption of CIQ-kmax is the largest because it keeps all the 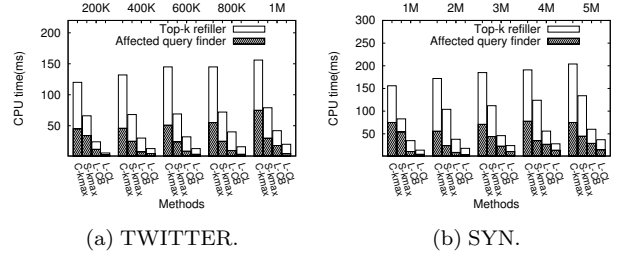queries in every leaf node of the quadtree. CCLS-CB spends the smallest amount of memory since it only maintains score bounds of cells for TR. SKYPE-kmax is the runner-up, followed by CCLS-CL. The gap between methods is not significant. All of them consumes several GB. Next we evaluate the effect of AQF and TR separately.

For AQF, we first compare the number of queries accessed, results shown in Figure 7a. CIQ accesses the most number of queries, followed by SKYPE. Thanks to l-signatures, CCLS accesses 25% to 34% queries compared to SKYPE. In Figure 7b, we plot the running times of AQF. Equipped with CC links and l-signatures, our method is faster than SKYPE by 2.3 to 2.6 times. For TR, we plot the number of objects accessed in Figure 8a. kmax performs the worst because it maintains a buffer of objects which updates frequently. Compared to kmax, CB's object access is moderately smaller. CL performs the best, accessing only 15% to 20% objects compared to kmax. The reduction in object access is converted to less running time in TR, as shown in Figure 8b. CB and CL are both faster than kmax. CL reaches a speed-up of 16.5 to 21.4 times over kmax.

9.4 Scalability

We study the scalability w.r.t. various parameters.

First, we vary the parameters of queries. Figure 9 plots the running times on YELP and TWITTER datasets by varying k from 10 to 50. We use C, S, L to denote CIQ, SKYPE, and CCLS in the figure, respectively. The running times are broken down to AQF and TR. As k increases, all the methods spend slightly more time in AQF and TR. This is because the k-th

object score decreases, and thus more cells, queries, and objects need to be accessed. We vary the number of keywords from 1 to 5 in queries and show the running times in Figure 10. The trend is similar to varying k, but the effect is more significant. In Figure 11, we vary α, the weight of spatial similarity, and plot the running times. When $\alpha = 1$ (or 0), only keyword (or spatial) similarity is used for ranking. The general trend is that the times decrease when we use a larger α. The reason is that a larger α indicates more weight on spatial information, which has better pruning power than keyword similarity.

Then we vary the parameters of objects. We adjust the ratio of location change to keyword change; e.g., 9:1 means location changes 9 times frequently than does keyword. This set of experiments was conducted on the synthetic dataset SYN as it is not supported by the real datsets. The running times are shown in Figure 12a. The general trend is that CIQ and SKYPE have their times increasing when keyword change is more frequent, while our methods slightly fluctuate. The main reason is that CIQ and SKYPE use single keywords to handle keyword similarity, and hence more sensitive to the frequency of keyword change. In contrast, our methods are equipped with l-signatures to reduce the time spent on processing keywords, thereby mitigating the influence of keyword change. In Figure 12b, we show the performance on SYN when varying the number of keywords changed per object update (we use fixed locations here). With more keywords changed in an update, all the methods fluctuate in running time but show a generally increasing trend. This is because more change of keywords leads to more change in the object's score, hence increasing the probability of more query and object access to compute top-k.

Next we vary the numbers of queries and objects as well as their ratio. Figures 13 and 14 show the running times when varying the numbers of queries and objects, respectively. We use TWITTER and SYN since they are larger than YELP. As expected, all the methods spend more time when we include more queries or objects. The effect of query number is more remarkable because we need to maintain k results and an object buffer (or a cell list) for one additional query, which accounts for more running time than having one more object. In Figure 15, we fix the sum of query and object numbers (2M on TWITTER and 5M on SYN) and vary their ratio (from 100:1 to 1:100). It can be seen that the running times decreases when there are more objects than queries. This is in accord with the results in Figures 13 and 14, showcasing more effect of query number than object number on query processing performance.

9.5 Batch Processing

We study the performance on batch processing by monitoring top-k for 100 batch updates, with batch size varying from 1 to 1,000. A batch size of 1 is exactly the case of a single object update. Figure 16 shows the average processing time **per object update** on YELP and TWITTER. We compare our methods with CIQ-kmax and SKYPE-kmax adapted for batch processing. Due to share computation of queries, all the methods exhibit a decreasing trend of average processing time for a larger batch size. Our methods substantially outperform alternative solutions (note the log-scale). In addition, the decrease rate of average processing time w.r.t. batch size is more remarkable for our methods, due to the one-signature-selection-for-all and $2k$-CL techniques tailored to batch processing. The average processing time of CCLS-CL for batch size 1,000 is 3.9ms on YELP and 4.3ms on TWITTER. Its speed-up over SKYPE-kmax is up to 14.9 and 20.0 times on the two datasets, respectively.

10. CONCLUSION

In this paper, we investigated a novel problem of top-k spatial keyword query processing on dynamic objects. We proposed an approach to monitoring top-k results for a large set of queries. Our approach consists of an affected query finder module and a top-k refiller module. We analyzed technical challenges and developed pruning strategies and corresponding data structures to speed up the two modules. Extension to batch processing was discussed. We conducted extensive experiments on real and synthetic datasets. The results showed that our solution outperforms alternative methods by 5 to 20 times speed-up.

11. REFERENCES

[1] P. Agrawal, A. Arasu, and R. Kaushik. On indexing error-tolerant set containment. In *SIGMOD*, pages 927–938, 2010.

[2] Amazon.com. Amazon prime air. `https://www.amazon.com/Amazon-Prime-Air/b?ie=UTF8&node=8037720011`, 2019.

[3] R. J. Bayardo, Y. Ma, and R. Srikant. Scaling up all pairs similarity search. In *WWW*, pages 131–140, 2007.

[4] L. Chen, G. Cong, and X. Cao. An efficient query indexing mechanism for filtering geo-textual data. In *SIGMOD*, pages 749–760, 2013.

[5] L. Chen, G. Cong, X. Cao, and K. Tan. Temporal spatial-keyword top-k publish/subscribe. In *ICDE*, pages 255–266, 2015.

[6] L. Chen, G. Cong, C. S. Jensen, and D. Wu. Spatial keyword query processing: An experimental evaluation. *PVLDB*, 6(3):217–228, 2013.

[7] L. Chen, S. Shang, Z. Zhang, X. Cao, C. S. Jensen, and P. Kalnis. Location-aware top-k term publish/subscribe. In *ICDE*, pages 749–760, 2018.

[8] G. Cong and C. S. Jensen. Querying geo-textual data: Spatial keyword queries and beyond. In *SIGMOD*, pages 2207–2212, 2016.

[9] G. Cong, C. S. Jensen, and D. Wu. Efficient retrieval of the top-k most relevant spatial web objects. *PVLDB*, 2(1):337–348, 2009.

[10] Y. Dong, H. Chen, and H. Kitagawa. Continuous search on dynamic spatial keyword objects. In *ICDE*, pages 1578–1581, 2019.

[11] Y. Dong, H. Chen, C. Xiao, K. Takeoka, M. Oyamada, J. X. Yu, and H. Kitagawa. Continuous top-k spatial-keyword search on dynamic objects. `https://dongyuyang.github.io/papers/dskfullversion.pdf`.

[12] C. Düntgen, T. Behr, and R. H. Güting. Berlinmod: a benchmark for moving object databases. *VLDB J.*, 18(6):1335–1368, 2009.

[13] L. Guo, J. Shao, H. H. Aung, and K. Tan. Efficient continuous top-k spatial keyword queries on road networks. *GeoInformatica*, 19(1):29–60, 2015.

[14] W. Huang, G. Li, K. Tan, and J. Feng. Efficient safe-region construction for moving top-k spatial keyword queries. In *CIKM*, pages 932–941, 2012.

[15] G. Li, Y. Wang, T. Wang, and J. Feng. Location-aware publish/subscribe. In *KDD*, pages 802–810, 2013.

[16] A. R. Mahmood and W. G. Aref. Query processing techniques for big spatial-keyword data. In *SIGMOD*, pages 1777–1782, 2017.

[17] K. Mouratidis, S. Bakiras, and D. Papadias. Continuous monitoring of top-k queries over sliding windows. In *SIGMOD*, pages 635–646, 2006.

[18] K. Mouratidis, M. Hadjieleftheriou, and D. Papadias. Conceptual partitioning: An efficient method for continuous nearest neighbor monitoring. In *SIGMOD*, pages 634–645, 2005.

[19] NYC Open Data. Green taxi trip data. `https://data.cityofnewyork.us/Transportation/2016-Green-Taxi-Trip-Data/hvrh-b6nb/`, 2016.

[20] Robomart. Self-driving stores. `https://robomart.co/`, 2019.

[21] J. B. Rocha-Junior, O. Gkorgkas, S. Jonassen, and K. Nørvåg. Efficient processing of top-k spatial keyword queries. In *SSTD*, pages 205–222, 2011.

[22] Twitter. Geotagged tweets from the us. `https://datorium.gesis.org/xmlui/handle/10.7802/1166`, 2016.

[23] P. Wang, C. Xiao, J. Qin, W. Wang, X. Zhang, and Y. Ishikawa. Local similarity search for unstructured text. In *SIGMOD*, pages 1991–2005, 2016.

[24] X. Wang, Y. Zhang, W. Zhang, X. Lin, and Z. Huang. SKYPE: top-k spatial-keyword publish/subscribe over sliding window. *PVLDB*, 9(7):588–599, 2016.

[25] X. Wang, Y. Zhang, W. Zhang, X. Lin, and W. Wang. Ap-tree: efficiently support location-aware publish/subscribe. *VLDB J.*, 24(6):823–848, 2015.

[26] D. Wu, M. L. Yiu, and C. S. Jensen. Moving spatial keyword queries: Formulation, methods, and analysis. *ACM Trans. Database Syst.*, 38(1):7:1–7:47, 2013.

[27] D. Wu, M. L. Yiu, C. S. Jensen, and G. Cong. Efficient continuously moving top-k spatial keyword query processing. In *ICDE*, pages 541–552, 2011.

[28] C. Xiao, W. Wang, X. Lin, J. X. Yu, and G. Wang. Efficient similarity joins for near-duplicate detection. *ACM Trans. Database Syst.*, 36(3):15:1–15:41.

[29] X. Xiong, M. F. Mokbel, and W. G. Aref. SEA-CNN: scalable processing of continuous k-nearest neighbor queries in spatio-temporal databases. In *ICDE*, pages 643–654, 2005.

[30] Yelp. Yelp dataset. `https://www.yelp.com/dataset`, 2019.

[31] K. Yi, H. Yu, J. Yang, G. Xia, and Y. Chen. Efficient maintenance of materialized top-k views. In *ICDE*, pages 189–200, 2003.

[32] N. E. Young. Greedy set-cover algorithms. In *Encyclopedia of Algorithms*. 2008.

[33] X. Yu, K. Q. Pu, and N. Koudas. Monitoring k-nearest neighbor queries over moving objects. In *ICDE*, pages 631–642, 2005.

[34] B. Zheng, K. Zheng, X. Xiao, H. Su, H. Yin, X. Zhou, and G. Li. Keyword-aware continuous knn query on road networks. In *ICDE*, pages 871–882, 2016.